

# Rohit Vemparala

[vemparalarohit1998@gmail.com](mailto:vemparalarohit1998@gmail.com)

| Boston, MA

| +1 (857) 540-0225

| [Github](#)

## SKILLS

**Languages:** Java, Python, Golang, Rust, TypeScript, Bash, C++, TLA+ (modeling concurrent/distributed systems)

**Frameworks:** Spring, React, Flask, NextJS, Netflix OSS (Eureka, Zuul, Hystrix)

**Cloud & Orchestration:** AWS (EKS, EC2, Lambda), Azure, Kubernetes, Docker (Swarm), Helm

**Data / Streaming:** Flink, Airflow, Kafka, Spark, Hadoop, Change Data Capture (CDC)

**DevOps & CI/CD:** GitLab CI/CD, Jenkins, ArgoCD, Snyk, SonarQube, Chaos Monkey

**Monitoring & Observability:** Datadog, Prometheus, OpenTelemetry, FluentBit, ELK, OpenSearch, PagerDuty, Zabbix

**Databases:** Relational (PostgreSQL, MySQL), NoSQL (DynamoDB, RocksDB, MongoDB), VectorDB (Milvus)

**Infrastructure as Code & Config Management:** Terraform, CloudFormation, AWS SAM, Ansible

**Certifications:** [Microsoft Certified: Azure Fundamentals](#)

## WORK EXPERIENCE

### Boston University

Research Fellow

Boston, MA

Jan 2025 - Present

- Built automated document parsing on air-gapped cluster using Donut (SWIN Transformer) and docTR
  - Fine-tuned models on **GPU cluster** - achieved **94%** accuracy on unseen documents

Directed Study/Teaching Assistant

Sept 2024 - Dec 2024

- Assisted in teaching the graduate level **Distributed Systems** course (GRS CS 651)
  - Worked on consensus algorithms (Raft, OmniPaxos) and modeling algorithm correctness with **TLA+**

### Staples Inc

Cloud Services Intern

Framingham, MA

June 2024 - Aug 2024

- Provide operational support for cloud Infrastructure - Azure Kubernetes Services upgrade, automate OS/security patching

### Protium

DevOps Engineer / Associate Vice President (Senior DevOps Engineer)

Bangalore, India

Jan 2022 - Aug 2023

- Managed multiple **Kubernetes** clusters across environments via **AWS CloudFormation** (AWS EKS)
  - Engine version upgrade with no downtime, setup multiple nodegroups (GPU, memory-intensive etc.)
  - Write **Helm** charts and **CI/CD** pipelines for deploying applications across stacks to clusters
- Led **cost optimization** effort saving over **\$10,000/month** on cloud infra spend
  - Optimize server costs (SPOT etc.) and placement to minimize data transfer cost (VPC Flow logs)
  - Minimize S3 - GET-api cost by merging files (parquet), storage tiering
- Handle **security and compliance** of applications and cloud infrastructure (CIS 1.4, AWS GuardDuty)
  - Integrated Gitlab **SAST**, **Snyk**, and image scanning into CI workflows to enforce secure coding
  - Set up Web Application Firewall (**WAF**), **AWS SSO** with **Azure Active Directory** (ABAC and RBAC)
  - Automate system patching with **Ansible**, setup server access via **Session Manager**
- Develop various applications, scripts, executable binaries with Python, Bash, Powershell
  - Cross platform executables (**Golang**) - providing cloud network access to external vendors
  - ETL** pipelines (**Apache Airflow**), **Flask** application integrated with **OpenSearch** for low-latency data retrieval

### Publicis Sapient

Associate Software Development Engineer

Bangalore, India

Jul 2020 - Dec 2021

- Built and maintain various systems for Goldman Sachs Transaction Banking platform
  - Rest APIs (**Spring Boot**, **React**), Kafka consumers, real-time dashboards (Reactor framework)
  - Stream processing (**Apache Flink** on **ECS Fargate**) millions of events in milliseconds from **Kafka** topics
  - Manage **RDS** database (Data archival, indexes, connection pooling, parameter tuning etc.)

## EDUCATION

### Boston University

MS in Computer Science (spec. Data-Centric Computing)

Sept 2023 - Jan 2025

GPA: 3.79/4.0

## PROJECTS

**Custom FS Syncer (S3-integrated File Tracker for K8s):** Built in **Go** with Linux **inotify**, syncs files to **S3** and supports rollback/versioning via **GitLab** CI. Designed for deployment in **Kubernetes** clusters

**Epidemic Engine (ETL & Anomaly Detection Pipeline):** Ingested health data from **Kafka**, processed via **Hadoop** batch jobs, and performed **Spark**-based anomaly detection on **Kubernetes**

**Pulse (Real-Time News Analysis Platform):** **RAG** system that ingests news with **WebSub**, generates insights with **LangChain** + **LLM**, streams a live knowledge graph into **Feldera**, an incremental query engine. Backend in **Rust**, **FastAPI** for LLM interface, Frontend in **Next.js** (Bun runtime, Tailwind), served via **Traefik** with **ngrok** tunnel, vector search with **Milvus**